Harini Reddy Edunuri

Fairfax, VA | harini.edunuri16@gmail.com | 571-663-7770 | linkedin.com/in/harinireddy07/
| https://github.com/hariniedunuri

PROFESSIONAL SUMMARY

- Data Analyst with 2+ years across Healthcare, Financial, and Retail analytics at B2B startups, building an AML/KYC pipeline that tiered 2,005 customers and flagged 290 AML-escalation candidates for compliance review through a Snowflake Unified Context Graph.
- Modeled chronic disease risk across all 50 U.S. states, surfacing socioeconomic and insurance-coverage disparities, by building 14 SQL, R, and Python visualizations deployed on AWS.
- Engineered a pricing analytics pipeline across 636 live-event listings, uncovering price spread as a 17x stronger predictor of resale value than artist demand, by comparing correlation r-values in Python and Power BI.

PROFESSIONAL EXPERIENCE

Accure — **Jan 2026 - Present**

Data Analyst — *Fairfax, VA*

Enterprise AI company whose AccureIQ platform turns regulated enterprise data into compliance and risk intelligence through private GenAI and a unified context graph.

- Tiered 2,005 customers into low, medium, and high risk bands and flagged 290 AML-escalation candidates, cutting manual AML review scope by 85% , by engineering risk features into a Snowflake Unified Context Graph in AccureIQ.
- Built a Snowflake warehouse of 8 interconnected tables across a 2,005-record KYC dataset, reducing data reconciliation time by 40% , by designing relational SQL schemas and validating integrity in Python and pandas.
- Generated explainable Customer 360 profiles of identity, document, and risk checks,cutting compliance-audit lookup time by 50% , by configuring an MCP server to link Snowflake to AccureIQ for real-time querying.
- Delivered the KYC intelligence platform across 5 Agile sprints on a 6-person cross-functional team, culminating in the final client showcase, by tracking tasks in YouTrack and managing version control in GitHub.

CustomerInsights.AI — **Jan 2023 - Aug 2024**

Data Analyst

Analytics service company for the life sciences industry, delivering commercial analytics to pharmaceutical and biotech clients through its proprietary ciPARTHENON AI platform.

- Built 15+ commercial analytics dashboards for 8+ life sciences clients on the proprietary ciPARTHENON platform, cutting report turnaround from 5 days to 2 for commercial teams , by cleaning and validating multi-source datasets for accuracy across all downstream visualizations.
- Streamlined stakeholder reporting by consolidating layered dashboard views into a single unified product, removing the need for end users to navigate separate views, by publishing finished dashboards via API for downstream systems to consume without manual handoffs.
- Standardized a repeatable dashboard-build workflow across 12+ client engagements, reducing rework and cutting delivery time by 30% , by cleaning, layering, consolidating, and API-publishing analytics in ciPARTHENON.

TECHNICAL PROJECTS

Chronic Disease Indicators Analysis | *github.com/hariniedunuri/chronic-disease-indicators-analysis*

SQL, Python (pandas, matplotlib), R (ggplot2), AWS (S3)

- Analyzed CDC Chronic Disease data across all 50 U.S. states, uncovering a positive correlation between income inequality and diabetes prevalence, by stratifying smoking and obesity in SQL and modeling income in R and ggplot2.
- Deployed the pipeline on AWS S3 and produced 14 visualizations across SQL, R, and Python, identifying the highest- and lowest-obesity states nationwide, by analyzing healthcare coverage against obesity in Python and pandas.

Duolingo Review Churn Analysis | *github.com/hariniedunuri/duolingo-review-churn-analysis*

Python (pandas, scikit-learn), SQL (SQLite), VADER, TF-IDF, Logistic Regression, Power BI

- Predicted churn risk at 0.91 ROC-AUC and 79% recall across 38,819 Google Play reviews, isolating ad frequency as the top driver of negative sentiment, by training a logistic regression classifier on TF-IDF features in scikit-learn.

- Built a 4-page Power BI dashboard of trend KPIs, complaint drivers, and a high-risk review explorer, surfacing a 3.08x concentration of ad complaints in low-rated reviews, by scraping and cleaning reviews through a custom Python ETL script.

Resale Ticket Pricing Analysis | *github.com/hariniedunuri/ticket-resale-pricing*
Python (pandas, seaborn), SQL (SQLite), Ticketmaster Discovery API, Power BI

- Delivered an interactive Power BI dashboard with 3 pricing recommendations for concert promoters and venues, projected $50K+ in recovered underpriced inventory , by building an analytics pipeline over 636 U.S. events across 16+ genres from the Ticketmaster Discovery API.
- Corrected a misleading pricing trend by validating mean against median across 858 snapshots, avoiding an outlier-driven false finding, by engineering a 6-window stratified sampling method in Python to remove date-clustering bias across 1,000+ events.

U.S. Airline Delay Forecasting | *github.com/hariniedunuri/airline-delay-analytics*
Python (pandas, NumPy, scikit-learn), statsmodels (SARIMA), seaborn, matplotlib

- Forecasted monthly airline delay hours across 398,233 U.S. flight records, revealing seasonal delay peaks in June-July and December for capacity planning, by tuning a SARIMA model in statsmodels to an MAE of ~17.3 on held-out data.
- Segmented airlines into low, moderate, and chronic high-delay tiers and flagged high-delay events at 99.38% accuracy, distinguishing chronic underperformers from reliable carriers, by training a logistic regression classifier and K-Means clustering on a 75/25 stratified split in scikit-learn.

TECHNICAL SKILLS & RELEVANT CERTIFICATIONS

Methodologies: Data Cleaning & Validation, ETL, Statistical Analysis, Predictive Modeling, Time-Series Forecasting, Feature Engineering, NLP (VADER, TF-IDF), Risk Classification, Data Visualization & Dashboarding, Stakeholder Reporting, Agile/Scrum

Tools: Power BI, Microsoft Excel (Pivot Tables, VLOOKUP), Snowflake, AWS (S3), MCP Server, pandas, NumPy, scikit-learn, statsmodels, matplotlib, seaborn, ggplot2, SQLite, Git, GitHub

Programming Languages: Python, SQL, R

Spoken Languages: English (Professional Working Proficiency), Telugu (Native)

Relevant Certifications: AWS Academy Graduate: Cloud Foundations , AWS Academy Graduate: Data Engineering

EDUCATION

George Mason University*, M.S., Data Analytics Engineering* — **Graduation Date: May 2026**
GPA: 4.00/4.0 | **Relevant Coursework:** Applied Machine Learning, Data Mining, Visual Analytics

SR University*, B.Tech, Electronics and Communication Engineering* — **Graduation Date: May 2024**
GPA: 3.60/4.0

Publications: Edunuri H, et al. "Predicting Real-Time House Prices: A Machine Learning Approach Using XGBoost Algorithm," 2024 IEEE Asia Pacific Conference on Innovation in Technology (APCIT), Mysuru, India. Achieved 91.77% prediction accuracy; 1 of 6 authors.